Module 5

Data I/O + Subset

Andrew Jaffe
Instructor

Data Output

While its nice to be able to read in a variety of data formats, it's equally important to be able to output data somewhere.

`write.table()`: prints its required argument `x` (after converting it to a `data.frame` if it is not one nor a `matrix`) to a file or connection.

```
write.table(x,file = "", append = FALSE, quote = TRUE, sep = " ",
            eol = "\n", na = "NA", dec = ".", row.names = TRUE,
            col.names = TRUE, qmethod = c("escape", "double"),
            fileEncoding = "")
```

Data Output

`x`: the R `data.frame` or `matrix` you want to write

`file`: the file name where you want to R object written. It can be an absolute path, or a filename (which writes the file to your working directory)

`sep`: what character separates the columns?

   "," = .csv - Note there is also a `write.csv()` function

   "\t" = tab delimited

`row.names`: I like setting this to FALSE because I email these to collaborators who open them in Excel

Data Output

For example, from the Homework 1 Dataset:

```
> dat = read.csv("data/Charm_City_Circulator_Ridership.csv", header=TRUE,as.is=TRUE)
> dat2 = dat[,c("day","date", "orangeAverage","purpleAverage","greenAverage",
+                "bannerAverage","daily")]
> write.csv(dat2, file="data/charmcitycirc_reduced.csv", row.names=FALSE)
```

Note that `row.names=TRUE` would make the first column contain the row names, here just the numbers `1:nrow(dat2)`, which is not very useful for Excel. Note that row names can be useful/informative in R if they contain information (but then they would just be a separate column).

Data Input - Excel

Many data analysts collaborate with researchers who use Excel to enter and curate their data. Often times, this is the input data for an analysis. You therefore have two options for getting this data into R:

Saving the Excel sheet as a .csv file, and using `read.csv()`

Using an add-on package called `xlsx`

For single worksheet .xlsx files, I often just save the spreadsheet as a .csv file (because I often have to strip off additional summary data from the columns)

For an .xlsx file with multiple well-formated worksheets, I use the `xlsx` package for reading in the data.

More on Packages

Packages are add-ons that are commonly written by users comprised of functions, data, and vignettes

Use `library()` or `require()` to load the package into memory so you can use its functions

Install packages using `install.packages("PackageName")`

Use `help(package="PackageName")` to see what contents the package has

http://cran.r-project.org/web/packages/available_packages_by_name.html

Some useful data input/output packages

foreign package - read data from Stata/SPSS/SAS

sas7bdat - read SAS data

xlsx - reads in XLS files

Installing Packages

```
> ## install.packages("xlsx",repos="http://cran.us.r-project.org")
> library(xlsx) # or require(xlsx)
```

Note you will need a stand-alone version of Java to use this

Saving R Data

It's very useful to be able to save collections of R objects for future analyses.

For example, if a task takes several hours(/days) to run, it might be nice to run it once and save the results for downstream analyses.

`save(...,file="[name].rda")`

where `...` is as many R objects, referenced by unquoted variable names, as you want to save.

For example, from the homework:

```
> save(dat,dat2,file="data/charmcirc.rda")
```

Saving R Data

You also probably have noticed the prompt when you close R about saving your workspace. The workspace is the collection of R objects and custom R functions in your current environment. You can check the workspace with `ls()` or view it in the "Workspace" tab:

```
> ls()
```

```
[1] "dat"   "dat2"  "Index" "tmp"    "x"
```

Saving the workspace will save all of these files in your current working directory as a hidden file called ".Rdata". The function `save.image()` also saves the entire workspace, but you can give your desired file name as an input (which is nicer because the file is not hidden).

Note that R Studio should be able to open any .rda or .Rdata file. Opening one of these file types from Windows Explorer or OSX's Finder loads all of the objects into your workspace and changes your working directory to wherever the file was located.

Loading R Data

You can easily load any '.rda' or '.Rdata' file with the `load()` function:

```
> tmp=load("data/charmcirc.rda")
> tmp
```

```
[1] "dat"  "dat2"
```

```
> ls()
```

```
[1] "dat"   "dat2"  "Index" "tmp"    "x"
```

Note that this saves the R object names as character strings in an object called 'tmp', which is nice if you already have a lot of items in your working directory, and/or you don't know exactly which got loaded in

Removing R Data

You can easily remove any R object(s) using the `rm()` or `remove()` functions, and they are no longer in your R environment (which you can confirm with running `ls()`)

You can also remove all of the objects you have added to your workplace with:

```
rm(list = ls())
```

Subsetting Data

Often you only want to look at subsets of a data set at any given time. As a review, elements of an R object are selected using the brackets.

Today we are going to look at more flexible ways of identifying which rows of a dataset to select.

Subsetting Data

You can put a – before integers inside brackets to remove these indices from the data.

```
> x = c(1,3,77,54,23,7,76,5)
> x[1:3] # first 3
```
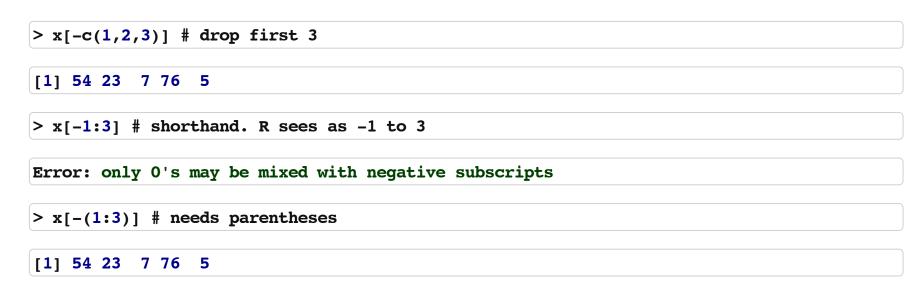
```
[1]   1   3 77
```

```
> x[-2] # all but the second
```

```
[1]   1 77 54 23   7 76   5
```

Subsetting Data

Note that you have to be careful with this syntax when dropping more than 1 element:

```
> x[-c(1,2,3)] # drop first 3
```

```
[1] 54 23  7 76  5
```

```
> x[-1:3] # shorthand. R sees as -1 to 3
```

```
Error: only 0's may be mixed with negative subscripts
```

```
> x[-(1:3)] # needs parentheses
```

```
[1] 54 23  7 76  5
```

Selecting on multiple queries

What about selecting rows based on the values of two variables? We can 'chain' together logical statements using the following:

    & : AND

    | : OR

```
> # which Mondays had more than 3000 average riders?
> which(dat$day =="Monday" & dat$daily > 3000)[1:20]
```

```
 [1] 148 155 162 169 176 183 190 197 204 211 218 225 232 239 246 253 260
[18] 267 274 281
```

AND

Which days had more than 10000 riders overall and more than 3000 riders on the purple line?

```
> Index=which(dat$daily > 10000 & dat$purpleAverage > 3000)
> length(Index) # the number of days
```

```
[1] 280
```

```
> head(dat[Index,],2) # first 2 rows
```

```
          day       date orangeAverage purpleAverage greenAverage
551    Friday 07/15/2011          4705          6293           NA
552 Saturday 07/16/2011          4624          7622           NA
    bannerAverage daily
551            NA 10998
552            NA 12246
```

OR

Which days had more than 10000 riders overall or more than 3000 riders on the purple line?

```
> Index=which(dat$daily > 10000 | dat$purpleAverage > 3000)
> length(Index) # the number of days
```

```
[1] 600
```

```
> head(dat[Index,],2) # first 2 rows
```

```
        day       date orangeAverage purpleAverage greenAverage
180    Friday 07/09/2010          2847          3094           NA
188 Saturday 07/17/2010          1513          3562           NA
    bannerAverage daily
180            NA  5941
188            NA  5076
```

Subsetting with missing data

Note that logical statements cannot evaluate missing values, and therefore returns an NA:

```
> dat$purpleAverage[1:10] > 0
```

```
 [1]  NA NA NA NA NA NA NA NA NA NA
```

```
> which(dat$purpleAverage > 0)[1:10]
```

```
 [1]  148 149 150 151 152 153 154 155 156 157
```

Subsetting with missing data

You can use the `complete.cases()` function on a data frame, matrix, or vector, which returns a logical vector indicating which cases are complete, i.e., they have no missing values.

Selecting on multiple queries

You can select rows where a value is allowed to be several categories. In the homework, we had to subset the Charm City Circulator dataset by each day. How can we select rows that are 1 of 2 days?

The `%in%` operator proves useful: "`%in%` is a more intuitive interface as a binary operator, which returns a logical vector indicating if there is a match or not for its left operand." It also returns `FALSE` for `NAs`

```
> (dat$day %in% c("Monday","Tuesday"))[1:20] # select entries that are monday or tuesday
```

```
 [1]   TRUE   TRUE FALSE FALSE FALSE FALSE FALSE   TRUE   TRUE FALSE FALSE
[12] FALSE FALSE FALSE   TRUE   TRUE FALSE FALSE FALSE FALSE
```

```
> which(dat$day %in% c("Monday","Tuesday"))[1:20] # which indices are true?
```

```
 [1]   1   2   8   9  15  16  22  23  29  30  36  37  43  44  50  51  57  58  64  65
```

Subsetting columns

We touched on this last class. You can select columns using the variable/column names or column index

```
> dat[1:3, c("purpleAverage","orangeAverage")]
```

```
  purpleAverage orangeAverage
1            NA           952
2            NA           796
3            NA          1212
```

```
> dat[1:3, c(7,5)]
```

```
  daily greenAverage
1   952           NA
2   796           NA
3  1212           NA
```

Subsetting columns

You can also remove a column by setting its value to NULL

```
> tmp = dat2
> tmp$daily=NULL
> tmp[1:3,]
```

```
        day       date orangeAverage purpleAverage greenAverage
1    Monday 01/11/2010           952            NA           NA
2   Tuesday 01/12/2010           796            NA           NA
3 Wednesday 01/13/2010          1212            NA           NA
  bannerAverage
1            NA
2            NA
3            NA
```